

Transactions and ACID properties

Introduction to Database Design 2011, Lecture 13



Course overview

- Communicating with DBMSs
- Designing databases
- Making databases efficient
- Making databases reliable



- Transactions
- ACID properties
- Concurrency and safe schedules



Transactions

- Transactions are units of work
- Example
 - Read balance of account A
 - Compute new balance = $\text{balance}(A) - 50$
 - Write new balance of A
 - Read balance of account B
 - Compute new balance = $\text{balance}(B) + 50$
 - Write balance of account B
- Transaction should be atomic: all or nothing



Transactions

- In database practice transactions are sequences of SQL statements
- Possibly intertwined by computations
- Can be written in
 - programming language (e.g. Java) accessing a database
 - procedural component of SQL
- `SQL:begin atomic ... end`
- Here we simplify and consider just sequences of reads and writes



- ACID stands for
 - Atomicity
 - Consistency
 - Isolation
 - Durability



Consistency

- Database must always be consistent wrt real world rules, e.g.,
 - Integrity constraints such as referential integrity must be satisfied
 - Rules of real world situation must be satisfied, e.g.,
 - Account balance must always be above a certain number (e.g. 0)
- Should also reflect real world as it is now
 - e.g. balance stored should correspond to actual balance of account



Transactional consistency

- Transaction leaves database in consistent state
 - ▀ (may assume database consistent before transaction start)
- May be required to satisfy other rules, e.g. leave the sum of the balances unchanged
 - ▀ (no money created or lost)
- During transaction consistency requirement may be violated temporarily
- Transactional consistency responsibility of transaction designer



Atomicity

- Transaction can be considered a unit of work
 - All or nothing!
- Consistency is impossible without atomicity
- Sometimes not possible to complete a started transaction, e.g.
 - In case of hardware failure or loss of connection
 - The application program may choose to abandon transaction
 - The DBMS may refuse to complete the transaction
- In these cases we say that transaction **fails**



Atomicity

- If a transaction fails it must be **aborted**
- This involves **rolling back** the transaction
- i.e., undoing all changes made by transaction
- Concurrency makes this complicated
 - e.g., changes made by transaction may have already been read
- Ability to roll back is implemented using a **log** of changes made to the database



Durability

- Durability is about trustworthy storage
- A transaction that has successfully completed is said to be **committed**
- Changes made by a committed transaction must be durable, i.e., able to survive
 - Power failure
 - Hardware failure etc




Durability

- When committing, changes must be written to non-volatile storage
- In practice, log is written to non-volatile storage
- Storage must be able to survive hardware failure
 - Maintain multiple copies of data
 - RAID



Transaction model



- Transactions may not interfere with each other
- In reality transactions are executed concurrently
- Statements of transactions intertwined
- DBMS should create illusion of transactions being executed sequentially
- Isolation is necessary for consistency



Need for concurrency

- Resources operate in parallel
 - Multiple CPUs
 - Data stored on multiple disks
- Computation may be stalled while waiting for data
- Gains of concurrency
 - Increased throughput: idle resources can be utilised
 - Decreased waiting time: new transactions can start executing immediately



Concurrency and safe schedules



A serial schedule

T_1	T_2
read(A)	
$A := A - 50$	
write(A)	
read(B)	
$B := B + 50$	
write(B)	
	read(A)
	$A := A + 20$
	write(A)



A good concurrent schedule

T_1	T_2
read(A) A := A - 50 write(A)	
	read(A) A := A + 20 write(A)
read(B) B := B + 50 write(B)	



A bad concurrent schedule

T_1	T_2
read(A) $A := A - 50$	
	read(A) $A := A + 20$ write(A)
write(A) read(B) $B := B + 50$ write(B)	



Schedules

- The DBMS receives a sequence of read and write requests from different transactions
- A schedule is an ordering of the reads and writes respecting the internal ordering in each transaction



Examples

- Two schedules

T_1	T_2	T_1	T_2
read(A)	read(A) write(A)	read(A) write(A)	read(A) write(A)
write(A) read(B) write(B)		read(B) write(B)	

- Schedule on right is equivalent to first T_1 then T_2
- It is up to DBMS to avoid bad schedules such as the one on the left



Conflicting operations

- Two operations **commute** if the order in which they are executed does not affect the result

$$\text{read}(A), \text{read}(B) = \text{read}(B), \text{read}(A)$$

$$\text{write}(A), \text{read}(B) = \text{read}(B), \text{write}(A)$$

$$\text{write}(A), \text{write}(B) = \text{write}(B), \text{write}(A)$$

$$\text{read}(A), \text{read}(A) = \text{read}(A), \text{read}(A)$$

- If they do not commute we say that they **conflict**

$$\text{write}(A), \text{read}(A) \neq \text{read}(A), \text{write}(A)$$

$$\text{write}(A), \text{write}(A) \neq \text{write}(A), \text{write}(A)$$



Conflict equivalence

- Two schedules are **conflict equivalent** if they differ only up to swapping commuting operations

T_1	T_2	T_1	T_2
read(A)		read(A)	
write(A)		write(A)	
	read(A)	read(B)	
	write(A)	write(B)	
read(B)			read(A)
write(B)			write(A)

- Executing conflict equivalent schedules gives same result



A non-example

- The following schedules are not conflict equivalent

T_1	T_2	T_1	T_2
read(A)			read(A)
write(A)			write(A)
	read(A)	read(A)	
	write(A)	write(A)	
read(B)		read(B)	
write(B)		write(B)	

- Suppose e.g. T_1 transfers all money available in A to B



Conflict serializability

- A schedule is **conflict serializable** if it is conflict equivalent to a serial schedule
- Example:

T_1	T_2
read(A) write(A)	
	read(A) write(A)
read(B) write(B)	



A non-serializable schedule

- The following is neither conflict equivalent to T_1T_2 nor T_2T_1

T_1	T_2
read(A)	read(A)
	write(A)
write(A)	
read(B)	
write(B)	



Serializable schedules

- Serializable schedules are the ‘good schedules’
- Parallel executions of transactions
- But still maintain illusion of serial execution
- DBMS should ensure that only serializable schedules occur
- This is usually done using locks




Detecting non-serializability

T_1	T_2	T_3	T_4
read(A)	write(A)	read(A) write(B) read(C)	
read(B)			write(C)



Precedence graph

- A cycle, so not conflict serializable



- **Theorem.** A schedule is conflict serializable if and only if its precedence graph is acyclic




Another example

T_1	T_2	T_3	T_4
read(A)	write(C)	write(B) read(D)	read(C) write(D)
write(B)			
	write(A)		



Precedence graph



- Equivalent serial schedules

$T_1T_2T_3T_4$

$T_1T_3T_2T_4$



Yet another example

T_1	T_2	T_3
	read(B)	
write(B)		
write(A)	write(A)	
		write(A)

- Not conflict serializable
- But result the same as running

$$T_2 T_1 T_3$$



View serializability

- Two schedules are **view equivalent** if
 - Corresponding reads in the two schedules always read same value
 - The changes made to the database are always the same
- A schedule is **view serializable** if it is view equivalent to a serial schedule
- Schedule on previous slide is view serializable



View serializability

- The following is not view serializable

T_1	T_2	T_3
write(B) read(A)	read(B) write(A)	 write(A)

- To see this need to check all serial combinations



View serializability

- We have two notions of serializability
- Conflict serializable schedules are also view serializable
- (because swapping commuting operations does not change behaviour of schedule)
- View serializable schedules need not be conflict serializable
- (see example a few slides back)



- ACID requirements for databases
 - Atomicity, consistency, isolation, durability
- Isolation is an illusion
 - In reality transactions are evaluated in parallel
- Two notions of good schedules
 - Conflict serializability
 - View serializability
- For exam you should be able to use these

